

Nearest Neighbour Method

Mingjun Zhong

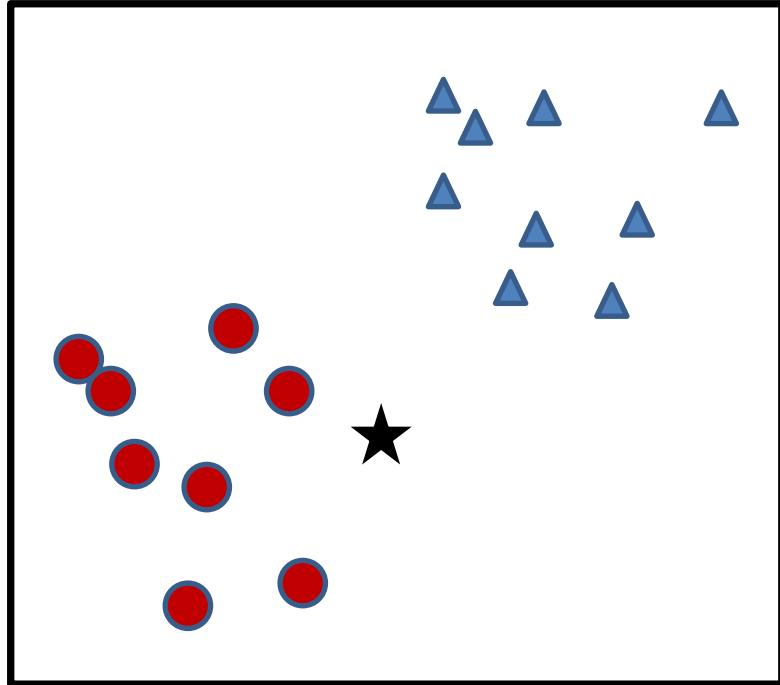
Department of Computing Science
University of Aberdeen

Outline

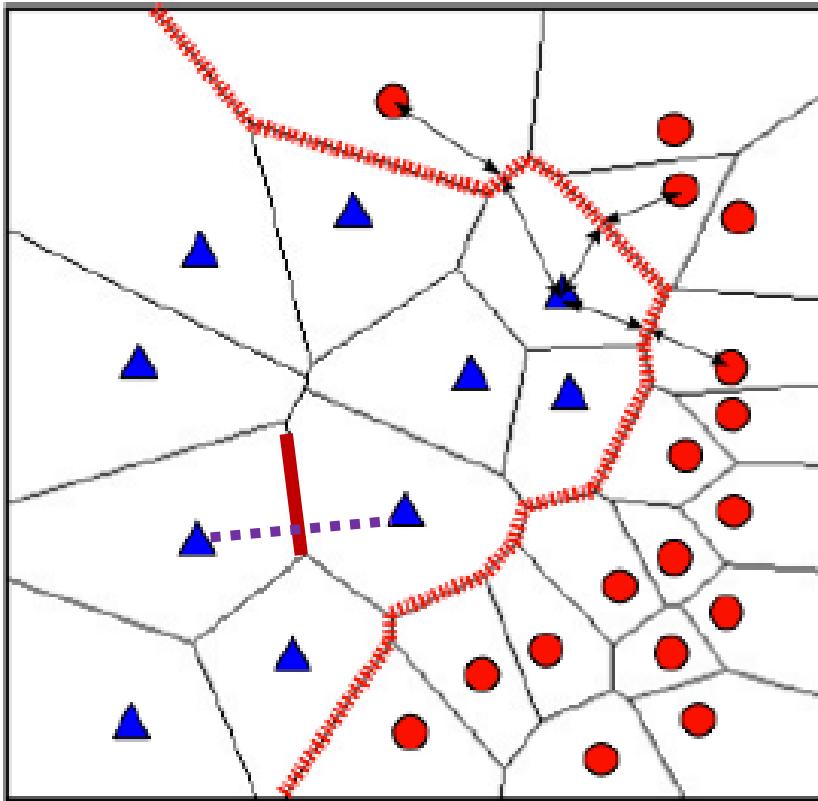
- Nearest neighbour method
 - Classification and regression
 - Practical issues: distance, k, ties, missing values
 - Optimality and assumptions
- Why should we normalise or standardise features?

Intuition for kNN

- Set of points (x, y) – two classes
- Which class is the star in?
- How would you do it?
- Nearby points are circles – use this as a basis for a learning algorithm



Nearest neighbour classification



- Use the intuition to classify a new point x :
 - Find the most similar training example x
 - Predict its class y
- Voronoi partition
 - Partitions space into regions
 - Boundary: points at same distance from two different training examples
- Classification boundary
 - Non-linear, reflects classes very well
 - Impressive for simple method

Nearest neighbour algorithm

➤ **Given:**

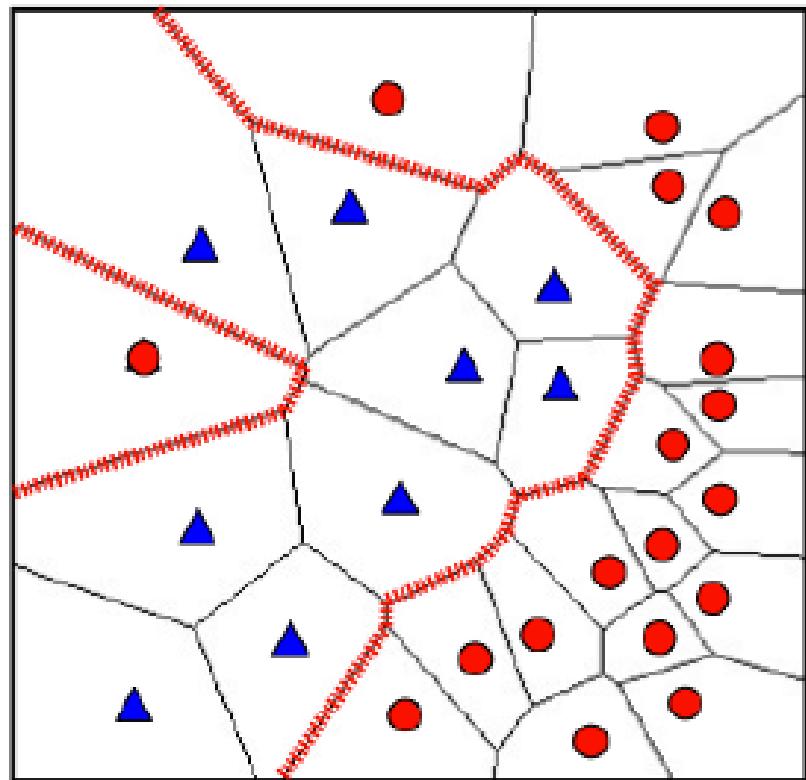
- Training examples $\{x_i, y_i\}$: x_i are attribute-value representation of examples; y_i class labels.
- Testing point x_* that we want to classify

➤ **Algorithm:**

- **Step 1:** compute distance from x_* to every training example x_i :
 $D(x_*, x_i)$.
- **Step 2:** select the closest instance(x', y')
- **Step 3:** output the class $y_* = y'$

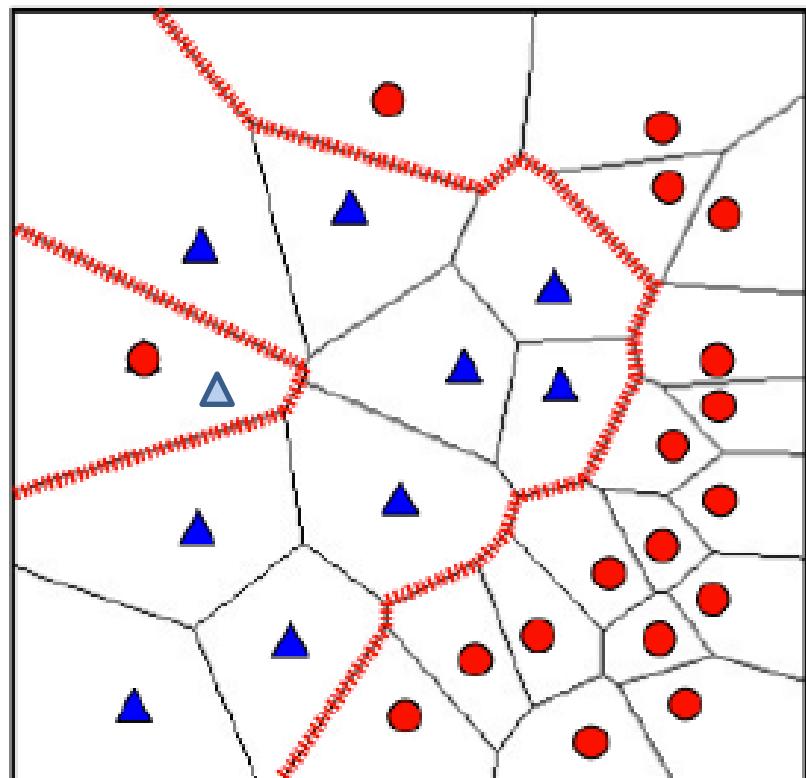
Nearest neighbour: outliers

- Algorithm is sensitive to outliers
 - Single mislabelled example dramatically changes boundary
- Idea:
 - Use more than one nearest neighbour to make decision
 - Count class labels in k most similar training examples



Nearest neighbour: outliers

- Algorithm is sensitive to outliers
 - Single mislabelled example dramatically changes boundary
- Idea:
 - Use more than one nearest neighbour to make decision
 - Count class labels in k most similar training examples



KNN classification algorithm

➤ **Given:**

- Training examples $\{x_i, y_i\}$: x_i are attribute-value representation of examples; y_i class labels.
- Testing point x_* that we want to classify

➤ **Algorithm:**

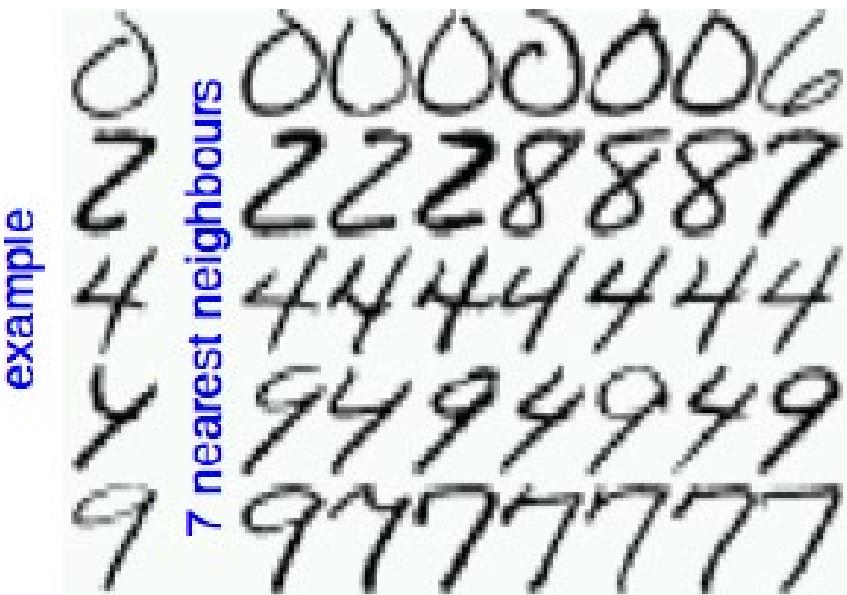
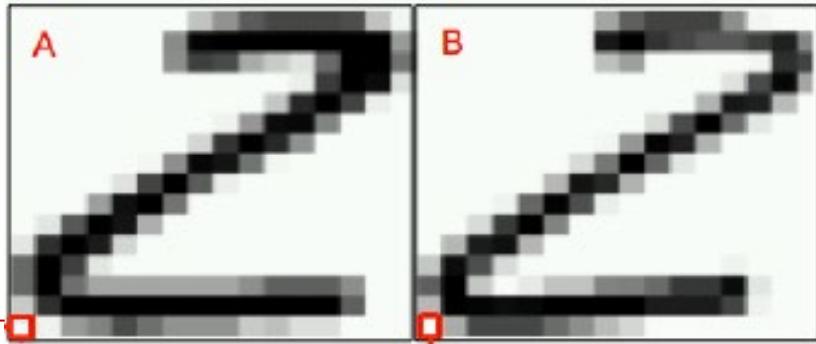
- **Step 1:** compute distance from x_* to every training example x_i : $D(x_*, x_i)$.
- **Step 2:** select k closest instances $(x_1, y_1), (x_2, y_2), \dots, (x_k, y_k)$
- **Step 3:** output the class y_* which is the most frequent class (the majority class) in $(x_1, y_1), (x_2, y_2), \dots, (x_k, y_k)$

Example: handwritten digits

- 16x16 grayscale bitmaps
- Euclidian distance:
 - On raw pixels

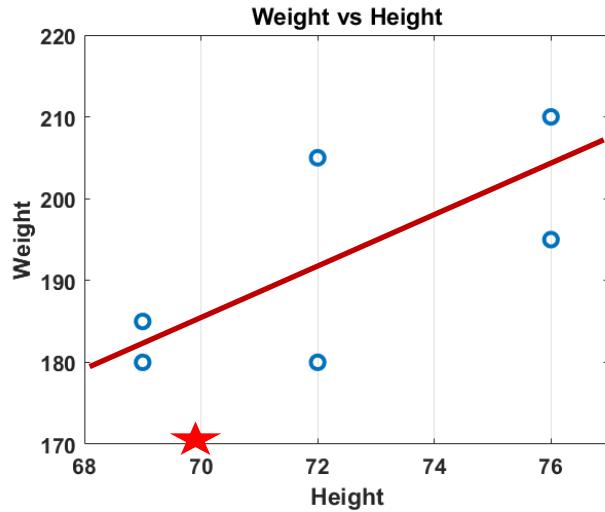
$$D(A, B) = \sqrt{\sum_r \sum_c (A_{r,c} - B_{r,c})^2}$$

- Accuracy:
 - K=7: 95.2%
 - Humans: 97.5%



KNN regression algorithm

- Given:
 - Training examples $\{x_i, y_i\}$: x_i are attribute-value representation of examples; y_i real-valued targets.
 - Testing point x_* that we want to predict the target
- Algorithm:
 - Step 1: compute distance from x_* to every training example x_i : $D(x, x_i)$.
 - Step 2: select k closest instances $(x_1, y_1), (x_2, y_2), \dots, (x_k, y_k)$
 - Step 3: output the target $y_* = \frac{1}{k} \sum_{j=1}^k y_j$



Example:

- Non-linear regression:

- $y = f(x)$

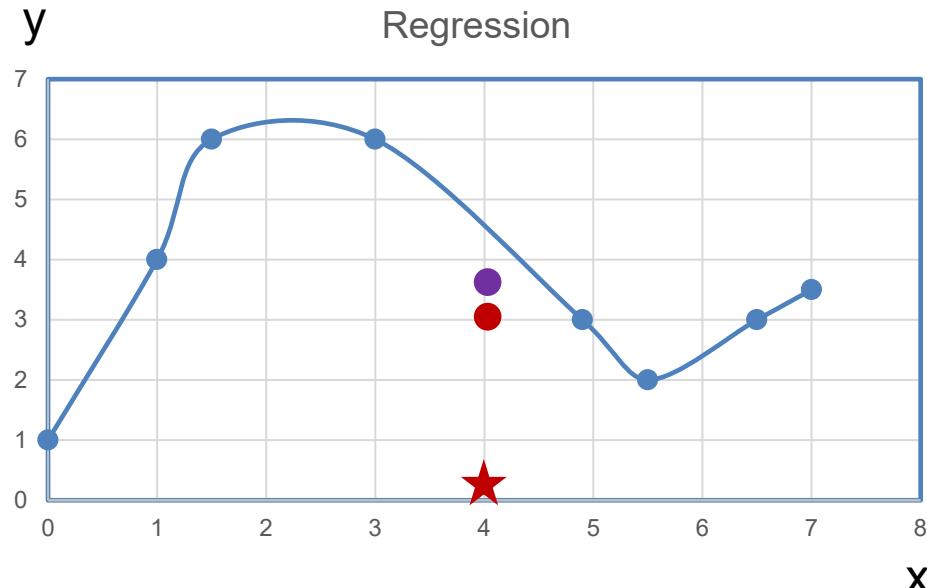
- Given:

Training Examples								
x	0	1	1.5	3	4.9	5.5	6.5	7
y	1	4	6	6	3	2	3	3.5

- Input $x = 4$, predict target y

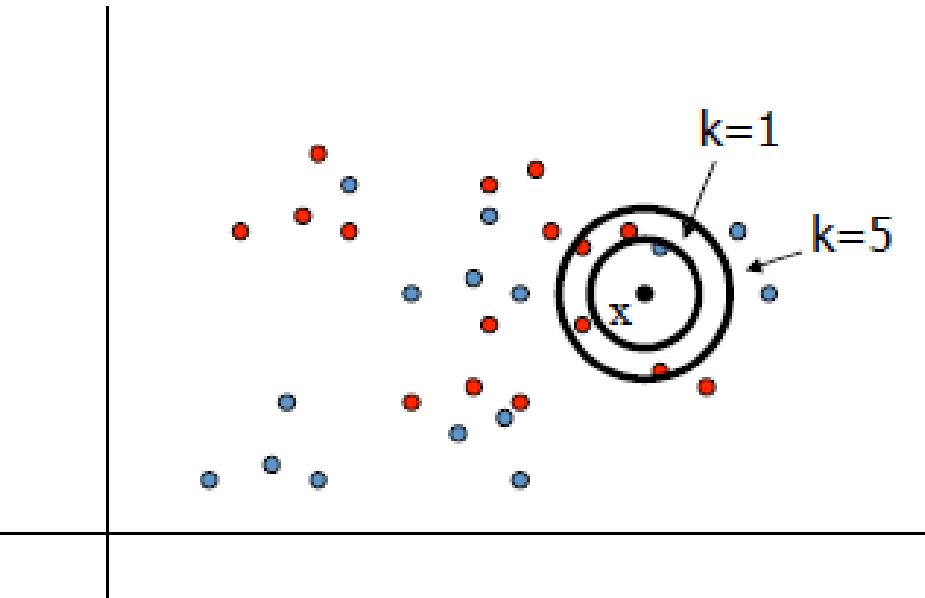
- 1-NN: $y = 3$ ●

- 3-NN: $y = \frac{11}{3}$ ●



Choosing the value of k

- Value of k has strong effect on kNN performance
 - Small value: unstable decision boundaries
 - Large value: everything classified to the most probable class
 - Affects smoothness of boundary
- Selecting the value of k (Cross Validation)
 - Set aside a portion of the training data (validation set)
 - Vary k , train model, validation error
 - Pick k which gives the best performance on validation set

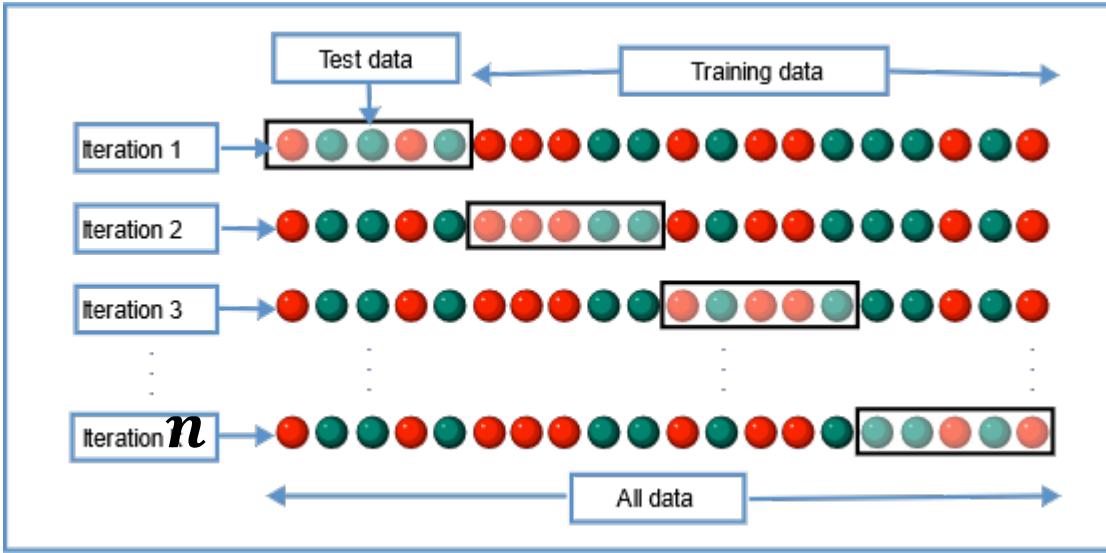


Choosing the value of k

- Using Cross-Validation method:
 - The k in KNN is a hyper-parameter, which influences the performance
 - Each k defines one KNN model
 - Cross-validation (CV) is a model selection method

Choosing the value of k

- What is cross-validation?
- Details to be introduced in later lectures
- n -fold cross-validation:
 - ✓ Shuffle the data samples
 - ✓ Equally split the data samples into n groups
 - ✓ For each k :
 - Iteratively use each group as the test data for validating the model, while using the rest data for training the model
 - In each iteration, compute the accuracy on the test data
 - Average the test accuracy over n iterations
 - ✓ Choose the k with the best average test accuracy.



Source: Wikipedia

Distance measures

- Key component of kNN
 - Defines which examples are similar; which are different
 - Has strong effect on performance

- Numeric attributes: Euclidian distance

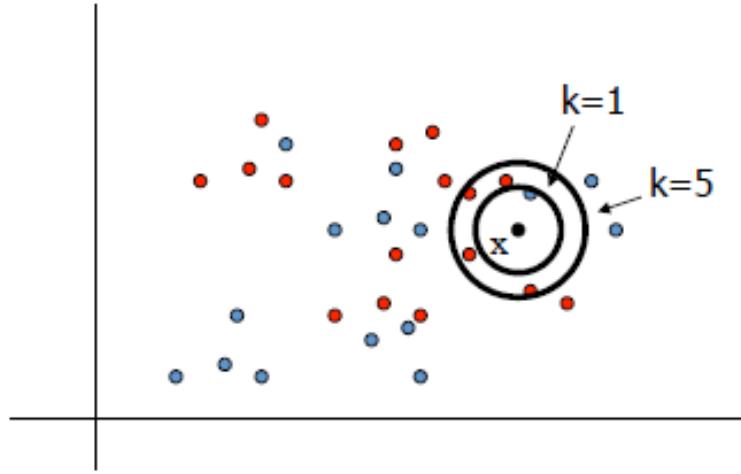
- $D(x, x') = \sqrt{\sum d(x_d, x'_d)^2}$

- Symmetric, spherical, treats all dimensions equally

- Categorical attributes: Hamming distance

- $D(x, x') = \sum_d \mathbf{1}_{x_d \neq x'_d}$

- Counts the number of different attributes



- $D(0110100, 1010101) = 3$
- $D(\text{"tube"}, \text{"cube"}) = 1$

Distance measures

➤ Minkowski distance:

- $D(x, x') = (\sum d(x_d, x'_d)^p)^{1/p}$

➤ Mahalanobis distance:

- $D(x, x') = \sqrt{(x - x')^T S^{-1} (x - x')}$

- S is the covariance matrix

➤ Kullback-Leibler (KL) divergence:

- $x_d > 0, \sum_d x_d = 1$

- $D(x, x') = -\sum_d x_d \log \frac{x_d}{x'_d}$

- $D(x', x) = -\sum_d x'_d \log \frac{x'_d}{x_d}$

- asymmetric

➤ Numeric attributes: Euclidian distance

- $D(x, x') = \sqrt{\sum d(x_d, x'_d)^2}$

- Symmetric, spherical, treats all dimensions equally

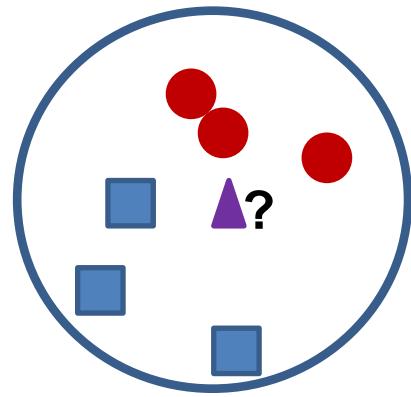
➤ Categorical attributes: Hamming distance

- $D(x, x') = \sum_d \mathbf{1}_{x_d \neq x'_d}$

- Counts the number of different attributes

Practical issues

- Resolving ties:
 - Equal number of positive and negative neighbours
 - Use odd k : does not solve multi-class
 - Breaking ties:
 - Random: flip a coin
 - Prior: pick the most numerous class
 - Nearest: use $1-nn$ to decide
- Missing values
 - Have to ‘fill in’, otherwise cannot compute distance
 - Reasonable choice: average value across entire dataset



Strengths and weaknesses

➤ Strengths

- Almost no assumptions about the data
- Smoothness: nearby regions of space assigned to same class
- Assumptions only made by distance function
- Non-parametric approach: “let the data speak for itself”
 - Nothing to infer from the data, except k and distance measure
 - Easy to update in online setting: just add new item to training set

➤ Weaknesses

- Need to handle missing values: fill-in or create a special distance
- Sensitive to class-outliers: mislabelled training instances
- Sensitive to lots of irrelevant attributes
- Computational expensive:
 - N: number of training examples; d: dimension of attributes
 - N grows: system will become slower and slower
 - Expensive in testing, not training time (too bad)

Summary

- Key idea: nearby points assigned to same class
- Can be used for classification and regression
- Simple, non-linear, non-parametric, asymptotically optimal
- Select k by cross validation
- Naïve implementations slow for big data

Normalising or Standardising Data Features

Normalisation or Standardisation

- Normalisation: scaling features to a range $[a, b]$, e.g., $[0, 1]$
 - Divide each feature value by the feature's norm value.
- Standardisation: make each feature have mean value = 0 and standard deviation = 1.
 - $x = (x_1, x_2, \dots, x_N)$: the data sample
 - $\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$: the mean value of x
 - $s = \sqrt{\frac{1}{N-1} \sum_{i=1}^N (x_i - \bar{x})^2}$: the standard deviation of x
 - These all can be calculated using Python

Normalisation or Standardisation

- Why is normalisation or standardisation important?
- Most of the models need to do such preprocessing before training:
 - Linear regression, logistic regression, k-nearest neighbour
- Why? :
 - Interpretation: making the feature units normalised
 - ✓ House sizes: **1300 and 1000**; numbers of bedrooms: **4 and 2**
 - ✓ Normalised or standardised features can interpret models
 - Make the learning process converge faster
 - ✓ Feature values which are not normalised make the Gradient Descent algorithm hard to converge (*have you tried last week's practical?*)
 - ✓ KNN algorithm: distance-based, if not normalised, distances are biased to those with large scales

$$\diamond \quad D(x, x') = \sqrt{\sum d(x_d, x'_d)^2} = \sqrt{(1300 - 1000)^2 + (4 - 2)^2}$$

❖ The effect of number of bedrooms is minor.